

Report

1. System's Perspective

1.1 Design and Architecture

When tasked with switching to another language for the system, C# was chosen. In correlation the project was built with Razor Pages and Entity Framework Core. The choice of C# was due to the groups already existing knowledge of the language, which streamlined the development process especially when it came to the architecture.

The architecture of the project follows a layered onion architecture split into the three parts: Core, Infrastructure and Web. The below visualization shows the responsibilities of each layer:

System Architecture

- The Core part of the system focuses on handling DTOs and repository interfaces. This layer does not reference any frameworks or libraries staying independent from the rest of the system.
- The Infrastructure focuses on the database context, migrations and the implementation of repository interfaces. This layer depends only on the Core, whilst having no reference to the Web layer.
- The Web handles the UI through razor pages along with the API. It also acts as the base of the system handling the dependency injection and referencing both the Core and Infrastructure layers of the application.

1.1.1 Choice of Final Infrastructure-as-Code Architecture We ended up migrating to Terraform towards the end of the project as it allows easy maintenance and resource control through defined interfaces. Terraform has a thoroughly defined Documentation for Digital Ocean resources, and the migration to defining existing Vagrant deployments along with "Click-Ops" resources therefore did not have much extra overhead.

1.2 Dependencies of MiniTwit

Technology	Stage	Role
Git / GitHub	Development, CI/CD	Source control, reviews, and work
Trello	Development, operations	Backlog management and work tr
Discord	Development, operations	Team communication and receivin
C# / .NET	Development, production	Application language and runtime
NuGet	Development, CI/CD	Package restore and feeds for .NE
ASP.NET Core (Razor Pages)	Development, production	Web UI and HTTP API
Entity Framework Core	Development, production	Database access and migrations
PostgreSQL	Development, testing, production	Primary data store (managed in I
Docker	Development, CI/CD, production	Container images and runtime iso

Technology	Stage	Role
Docker Hub	CI/CD, production	Registry for built application images
Docker Compose	Development, testing	Local and test multi-container setups
Docker Swarm	Production	Orchestration and rolling updates
DigitalOcean	Infrastructure, production	Cloud VMs, managed database, and DNS
DigitalOcean Spaces (S3-compatible)	Infrastructure, CI/CD	Object storage with an S3-compatible API
Terraform	Infrastructure, CI/CD	Infrastructure as code for managing cloud resources
GitHub Actions	CI/CD	Continuous integration and deployment
Third-party GitHub Actions	CI/CD	Marketplace and vendor-maintained actions
GitHub CLI	CI/CD	Command-line GitHub integration
Ubuntu	CI/CD, production	Base operating system on GitHub Actions
SSH (OpenSSH)	CI/CD, production	Remote deploy, file copy, and service control
Prometheus	Monitoring	Metrics collection
Grafana	Monitoring	Dashboards and alerting
Loki	Monitoring	Centralized log storage
Promtail	Monitoring	Shipping container logs to Loki
Python	Testing (local and CI/CD)	API simulator test driver and test runner
Selenium	Testing (local and CI/CD)	Browser-based UI tests (with ChromeDriver)
dotnet format	Development, CI/CD	C# formatting and verifying the code
Roslynator	Development, CI/CD	C# static analysis (Roslyn-based)
Codespell	Development, CI/CD	Spell checking across the repository
Hadolint	Development, CI/CD	Linting Dockerfiles
Codacy	Development, CI/CD	Hosted static analysis and pull-request checks
SonarCloud	Development, CI/CD	SonarQube-family analysis and quality gates
CodeQL	Development, CI/CD	Semantic security and quality scanning
Docker Scout	CI/CD	Container image vulnerability scanning
OpenAPI Generator	Development	Generating the API simulator stubs
Pandoc / LaTeX	CI/CD	Building the report PDF in automatic
GNU Make	Development, CI/CD	Task automation (local and in workflow)

Libraries NuGet references come from the Razor Pages solution (`razor-pages/Web`, `razor-pages/Infrastructure`). Direct Python libraries for automated tests are listed below; the UI test lockfile also pins transitive versions (see `tests/selenium/requirements.txt`).

Library	Context	Usage
Microsoft.AspNetCore.Identity.EntityFrameworkCore	Web, Infrastructure	ASP.NET Core Identity integration
Microsoft.EntityFrameworkCore.Design	Web, Infrastructure	EF Core design-time support and scaffolding
Npgsql.EntityFrameworkCore.PostgreSQL	Web, Infrastructure	EF Core provider for PostgreSQL
Newtonsoft.Json	Web	JSON serialization and deserialization
Swashbuckle.AspNetCore.Annotations	Web	OpenAPI metadata and attributes
Swashbuckle.AspNetCore.Newtonsoft	Web	OpenAPI generation using Newtonsoft.Json
DotNetEnv	Web	Loading environment variables

Library	Context	Usage
prometheus-net.AspNetCore	Web	Prometheus metrics for ASP.NET
Microsoft.CodeAnalysis.Analyzers	Infrastructure	Build-time Roslyn analyzers
Microsoft.Extensions.Hosting	Infrastructure	Hosting abstractions for background
prometheus-net	Infrastructure	Prometheus metric registration
Npgsql	Infrastructure	PostgreSQL data provider (ADO.NET)
TimeZoneConverter	Infrastructure	Resolving time zones in infrastructure
requests	tests/API_Spec	HTTP calls from the API simulator
pytest	tests/selenium	Test runner for the Selenium UI tests
selenium	tests/selenium	WebDriver client driving the browser

1.3 Current State of Our Systems

The current state of our system leaves it steadily functional across performance, scalability, code quality, security, and testing. However there are still some limiting factors that exclude it from being entirely production-ready.

In regards to performance and scalability, the primary issue we encountered is the resources available to the server being insufficient for exceptionally high traffic. However, that is relatively easily managed by upgrading the server plan on DigitalOcean.

The test coverage is quite extensive across the API and browser-based UI levels, but there could be more explicit tests for base application logic, as well as error and edge-case interactions and security behavior.

Static Analysis and Code Quality Tools

- **SonarQube** indicate some potential security hotspots and issues with reliability, but overall still rates the application A-ratings in security and maintainability.
- **Codacy** gives the application as a whole an A-rating, but also indicates a few potential security hotspots.
- **CodeQL** passes on all its vulnerability checks. It only fails on javascript due to a file being empty.
- **Hadolint** and **Roslynator** show no issues.

SonarQube's Quality Assessment: SonarQubeQuality

2. Process' perspective

2.1 CI/CD pipelines, deployment, and release

All development work is done on branches and requires a pull request to be merged into the main branch. Pull Requests are automatically checked with code scanning tools and also triggers a QA build which runs a full build, test

and deployment to a QA droplet and database. Note, that due to limitations on number of allowed droplets in our Digital Ocean account level, this QA droplet was later included in the production swarm as well.

After merging a pull request into main, the report pdf is built iff changes have been made in the relevant files. Nothing is immediately pushed to production as we deemed that we wanted our releases to contain more than a single small change and have more control of when releases to production were made. The control of timing is important to ensure stability of the application and timely action given a failure/bug.

We used an automated deployment pipeline to deploy our production services that automatically triggers when a tag is pushed to the repository. We attempted to follow a form of semantic versioning for tag names, to have a consistent format and a notion of how big each release was. The automated deployment builds a docker image and deploys the stack on the swarm leader node.

Monitoring is deployed manually in a separate workflow. The monitoring droplet was initially a stand-alone droplet, but given the Digital Ocean limitation on droplets, this droplet was also later included in the swarm. The monitoring deployment could have been automatically deployed if changes appeared in the relevant root folder, yet changes to the configurations were rather rare and we therefore did not find it necessary.

Below is an overview of the different stages of development towards operationalization. In the following sections we will deep dive into the QA deployment workflow, continuous deployment release workflow and the monitoring deployment workflow.

flowchart LR

```
subgraph dev [Development Branch]
    A[Branch Work] --> B[Make a PR]
    B --> C[QA Deployment workflow]
    B --> K[Automated Checks]
    C --> L[Peer Review]
    K --> L[Peer Review]
end
subgraph merge [Main Branch]
    L --> D[Merge to main]
    D --> E[Report build on report changes]
end
subgraph release [Release Tagging]
    D --> F[Semantic version tag]
    F --> G[Continuous Deployment]
    G --> H[Docker Swarm production]
end
subgraph ops [Operationalization]
    D --> I[Deploy Monitoring manual]
end
```

```

    I --> J[Prometheus / Grafana / Loki]
    H --> J
end

```

Pull-request pipeline (QA Deployment) The QA deployment is defined in `.github/workflows/continous-QA-deployment.yaml` and is automatically run on pull requests towards the main branch.

```

sequenceDiagram
    participant GH as GitHub Actions
    participant Hub as Docker Hub
    participant QA as QA droplet
    participant DB as PostgreSQL QA DB

    GH->>Hub: build and push testminitwit
    GH->>GH: terraform plan + PR comment
    GH->>QA: scp env + compose-test.yaml
    GH->>GH: Docker Scout CVE scan
    GH->>QA: Deploy Docker Compose
    Hub->>QA: Pull Latest Image
    QA->>QA: Docker Compose Up
    GH->>DB: reset schema / simulator state
    GH->>GH: run-all-validations
    GH->>QA: API simulator + Selenium against :8081
    QA->>DB: Query
    DB->>QA: Data
    QA->>GH: API / UI Results

```

Above flow chart shows the various steps and interactions between systems happening during the QA Deployment and test workflow. The workflow runs at the same time as the static code analysis tools `CodeQL`, `SonarCube` and `Codacy`.

Production release (Continuous Deployment) The Continuous deployment to production is defined in `.github/workflows/continous-deployment.yaml` and is automatically run on tags pushed to the main branch.

```

sequenceDiagram
    participant GH as GitHub Actions
    participant Hub as Docker Hub
    participant DO as DigitalOcean
    participant Mgr as Swarm manager
    participant Nodes as Swarm cluster
    participant DB as Managed PostgreSQL

    GH->>Hub: build and push minitwitimage:SHA
    GH->>DO: terraform apply prod
    Note over DO,DB: Applies IaC for VMs network and managed Postgres etc.

```

```

GH->>Mgr: scp compose.yaml
GH->>Mgr: docker stack deploy minitwit
Hub->>Nodes: pull image (with-registry-auth)
Mgr->>Nodes: rolling update (3 replicas)

```

Monitoring deployment The monitoring stack deployment is defined in `.github/workflows/monitor-deployment.yaml` and runs only when someone manually triggers it.

sequenceDiagram

```

participant GH as GitHub Actions
participant Hub as Docker Hub
participant Mon as Monitoring server
participant Mgr as Swarm manager

```

```

GH->>Mon: scp repo monitoring/* to /vagrant/monitoring
GH->>Mgr: scp repo monitoring/* to /vagrant/monitoring
GH->>Mgr: docker stack deploy monitoring
Mgr->>Mon: deploy constrained to server
Hub->>Mon: Pull Prometheus, Grafana and Loki Images
Mon->>Mon: Run Prometheus, Grafana and Loki Containers

```

Deployment & Release Summary

Environment	How it is updated	Image / orchestration
Local	<code>make app-build</code> (<code>compose-test.yaml</code> , port 8081)	Local Docker Compose
QA (pre-merge)	QA Deployment workflow on PR	<code>testminitwit:latest</code> , Compose on t
Production	Tag → Continuous Deployment	<code>minitwitimage:<sha></code> , Docker Swarm
Monitoring	Manual Deploy Monitoring workflow	Swarm stack <code>monitoring</code>

2.2 Monitoring

The monitoring of our application is done through the use of Prometheus and Grafana.

The data collection is handled by Prometheus' .NET client library, `prometheus-net`, with `UseMetricServer` collecting and exposing metrics for use by Grafana. Additional http request metrics are collected through the use of the `UseHttpMetrics` middleware provided by Prometheus. Custom metric gatherers were also implemented to retrieve metrics from the application's database.

Grafana is then used to retrieve these exposed metrics provided by Prometheus and allows for the construction of various visualizations. We also implemented a Grafana alert based on the up-time metric to inform us when the server was down.

Monitoring Panels

- Current and uptime of container status
- CPU utilization
- Memory utilized by dotnet processes, as well as total physical allocated memory
- Total tweets and tweet rate
- Total users and registration rate
- Http request response latency by their action
- Http GET and POST request rates over time by their response status codes

2.3 Aggregated logs

All assignment completions for each week have been aggregated in View project log . It was standard practice for everyone to document which tasks they completed. A type of "Meta" log used is the README file it serves as how we ought to implement the assignments as well as principle on how work as a group Docker has a built-in log system for each droplet. This logging system was rarely used except for some debugging cases. all live logs are shipped to grafana alt text

2.4 Security Hardening

The firewall for the DigitalOcean Droplets was configured to improve security. Inbound firewall rules were configured in DigitalOcean to provide an clear overview of the rules for each Droplet instead of `ufw`, such that the rules can be conveniently managed by each group member through the user interface. Inbound rules defines the traffic allowed to the Droplets on which ports and from which sources, and all other incoming traffic is blocked. In addition, Docker does not bypass DigitalOcean's firewalls. This ensures that the sensitive ports is only accessed by internal Droplets, preventing security risks from unnecessary exposed ports.

The production application were given firewall rules for:

- Standard internet and access ports (TCP 22, TCP 80, TCP 443)
- Docker port (TCP 2376)
- Docker Swarm infrastructure ports (TCP 2377, UDP 4789, TCP/UDP 7946)
- Application ports (TCP 8080)
- Grafana Loki log aggregation port (TCP 3100)
- Prometheus ports (TCP 9090, TCP 9095, TCP 9096, TCP 9100)

For the security hardening of our system a security assessment was made showing an overview of assets/threats/risks:

Assets

- Web Application & API Endpoint
- PostgreSQL Database
- Monitoring (Grafana, Loki, Promtail & Prometheus)

Threat Sources

- SQL Injection
- Cross-Site Scripting (XSS)
- DDoS Attack
- Brute force attack

Risk Scenarios

- An attacker uses SQL injection on the web application to gain access to the database.
- An attacker injects a script into a post, which is executed on other users' browsers, allowing the attacker to steal session tokens.
- An attacker uses a script to spam an API endpoint, resulting in a denial of service.

Risk Analysis				
Scenario	Likelihood	Impact	Risk	
SQL Injection	High/Common	High	Critical	
Cross-Site Scripting (XSS)	High/Common	High	High	
DDoS Attack	Medium/Uncommon	Medium	Medium	

To solve the various risk scenarios the following measures were taken:

- **SQL Injection:** All inputs are sanitized to avoid script injection. This is handled automatically by Entity Framework Core.
- **Cross-Site Scripting (XSS):** This gains from the input sanitation, but still needs an output encoding to ensure data is rendered as text. Which is handled by Razor pages rendering all posts as plain text.
- **DDoS Attack:** Access is restricted to only allow a certain amount of requests per/minute, to minimize the effect of DDoS attacks.

Other Security Measures

Other security measures were also taken such as:

- Setting up a firewall (using `ufw`) on the server only allowing specific traffic through on specified ports.
- Ensuring the application runs on HTTPS with a TLS certificate and setting up `Nginx` for a reverse proxy in front of the application.
- Docker images were also security hardened by ensuring only user privileges

2.5 Availability and Scaling

Availability and scaling in the MiniTwit application are managed by Docker Swarm. A Swarm cluster of the DigitalOcean Droplets is joined into a single Swarm cluster, which continuously monitors and enforces the declared desired state.

High availability is handled by having manager redundancy, three production container replicas, and automatic self-healing. All three nodes in the cluster are given the **manager** role to prevent a single point of failure if one of the manager nodes crashes. When a node in the cluster crashes, the Swarm detects a difference between the actual state and the declared desired state, such that the number of actual running replicas is lower than three, which triggers *self-healing* to restore the third replica.

Scaling is handled by Docker Swarm's built-in Ingress Routing Mesh which functions as a load balancer. Swarm evenly distributes incoming user requests across all three healthy replicas of the production container to handle high amounts of concurrent requests. This parallelizes the workload across the nodes, such that a container does not consume all the resources of a single node.

To ensure low downtime during the transition from the standalone containers to a Docker Swarm cluster, the stack was deployed with a **blue-green service deployment** strategy from the start, such that the running containers were gradually replaced by the updated ones. This is achieved by configuring the deployment settings within the Docker Compose file to set the update order to **start-first** and a delay parameter that dictates how long Docker Swarm should wait after starting an updated container before terminating an old one, allowing the new container to initialize.

As a result, the services remain available during deployment. By default, Docker Swarm uses the rolling update strategy which terminates the old container before starting a new one. This is called **stop-first**. By terminating the containers first, the default strategy forces the application to experience downtime during the window between container termination and container initialization.

While our strategy should have ensured low downtime during the transition to Docker Swarm, the production application still experienced downtime due to overlooked human errors. These errors came as a result of debugging separate Docker Swarm migration issues. Specifically, the Swarm Ingress routing mesh overwrote the port of the development environment on one of our DigitalOcean Droplets, and the Loki logs failed to display on our Grafana dashboards. During the debugging process, accidental downtime of the application was introduced when pulling the wrong container image due to misconfigured environment secrets, or changing the application to use another port, which prevented the simulator from reaching it.

To avoid these issues in the future, a solution could be to replicate the Docker Swarm infrastructure within an isolated development environment, so any configuration changes during the transition does not affect live production.

3. Reflection Perspective

3.1 Evolution and Refactoring

On first refactoring from Python to RazorPages with C# we ran into unforeseen issues with the methods not working as intended. This slowed us down but once the bugs were solved we were able to make our release. We had no issues Refactoring to our Onion Structure. It was time-consuming but that half the group being familiar with the framework made the process smooth.

We discussed that it may have been useful to have defined the infrastructure in Terraform from the beginning and that it may have led us to avoid having the amount of "Click-Ops" we had during the project (setting up a managed database, modifying network rules for droplets etc.) giving better reproducibility and version history.

3.2 Operation

To increase robustness we added a QA deployment on Pull Requests, requiring the application to be built, pushed and tests to pass before merging it into the main branch. This allowed us to test all of our features fully before releasing them to our main application and thereby decreased the amount of bugs and operational work.

In early April we started receiving warnings from the built-in resource alert system in Digital Ocean that our Database Cluster was above 90% CPU utilization. We started investigating the issue and realized that the amount of requests coming in had ramped up so much that our database could not follow along. We chose to resize the cluster such that it had an extra virtual CPU after cost-benefit analysis concluding that the developer time it would take to improve the ORM system to send fewer requests would be too time consuming versus the cost of upgrading the database cluster. The database was resized with no downtime.

Once we had fully migrated to swarm, including log shipping from all our droplets, the droplet containing the monitoring application ended up being overloaded such that our monitoring application became unreachable. This warranted an upgrade of the droplet containing the monitoring application. If we had access to spin up more droplets, we may have considered horizontal scaling instead of vertical. The monitoring droplet was resized using terraform and therefore had minimal possible downtime. The full process took ~6 minutes:

DigitalOcean monitoring droplet resize (duration ~6 minutes)

We struggled with the continuous deployment of our monitoring application in terms of ensuring new dashboards would appear and an unintended addition of a flag that reset the volumes for Loki and Prometheus. After realizing the issue and looking at a few different combinations of flags, we fixed the issues and accepted the loss of earlier metrics & logs.

An improvement of the monitoring deployment would be to trigger it on changes

to the particular folder containing monitoring definitions instead of only manual trigger.

3.3 Maintenance

4. Use of Generative AI

The generative AI tool Cursor was used to discuss issues and warnings throughout the project and provide guidance on issues that we had not encountered before. This was beneficial to the development process as it unblocked us in completing tasks. We used Cursor to improve the code in terms of maintainability by standardizing the format and structure along with industry standard formatting / linting tools (i.e. using a mix of CLI tools and GenAI). We also employed Cursor to summarize branch work in our log.md and other documentation, to ensure chore tasks were being done rather than neglected. These logs and documentation are used to remind us of the work and considerations throughout the different stages of the project which have been written / read for internal use.

Claude and ChatGPT were used to quickly bounce ideas off of and to help identify the pros/cons of options when many were presented. Additionally, they were very useful in quickly parsing large error logs.

ChatGPT was used for debugging GitHub Actions, Docker, and DevOps setup issues, it often hindered our work on the Deployment workflows but helped significantly on Command lines in the terminal.

The AI tool Github Copilot was used throughout the course for better code understanding, serving as interactive documentation. It was also used for very specific code fixes, usually hinted by SonarQube, such as `b6f71cf3242a25a0c03cd0c0763040417532838f - add wheel hashes to the requirements file`. This made it possible to implement security and maintainability solutions which sometimes felt like an overkill.

When it came to the use of GAI, ChatGPT and GitHub Copilot was used to better understand coding errors and thereby helping in solving them. Another aspect of using AI, would be on how to do specific things in different languages. For example "How do i write inline code in .md?" or "How do i change the rejection status code on the rate limiter?".

The **Google Gemini 3 model** has primarily been used for debugging and to resolve technical uncertainty during development.

This includes asking the generative AI model to:

- find possible fixes for errors, as well as explaining how and why they appeared.
- provide an overview and understanding of important features and commands of new tools and technologies.

- review developer decisions to ensure that any changes to the application is correct.
- check for grammar and phrasings when writing documentation or report.